

Oslo

OPEN SPEECH LANGUAGE OPTIMIZER

Universal, High-Performance Real-Time Speech Translation

Suraj Kunwar (1CR25AI120), Suryakanth(1CR25AI121), Prajwal
Santoshkumar Govi(1CR25IS114)

DEPT OF BASIC SCIENCES, CMRIT, BANGALORE

TABLE OF CONTENTS

CHAPTER 1: Introduction

CHAPTER 2: Problem Statement

CHAPTER 3: Project Objectives

CHAPTER 4: System Architecture

CHAPTER 5: Experimental Performance and Evaluation Benchmarks

CHAPTER 6: Conclusion and Future Roadmap

REFERENCES

Chapter 1

INTRODUCTION

In an increasingly globalized world, real-time cross-lingual communication has evolved from a luxury to a critical operational requirement across industries such as international business, customer support, emergency services, and live broadcasting. Conventional translation systems typically rely on a cascading architecture: a Voice Activity Detection (VAD) system isolates audio, an Automatic Speech Recognition (ASR) engine transcribes speech to text, and a Machine Translation (MT) framework processes the text into the target language.

While these technologies have matured individually, integrating them into a unified, synchronous, real-time pipeline introduces severe technical friction. In live environments, user experience dictates that end-to-end latency must remain near-conversational (typically under 1.5 seconds). However, standard implementations suffer from structural delays caused by sequential processing bottlenecks, unoptimized audio streams, unmanaged API round-trip constraints, and high computation overhead. Minimizing this latency while maintaining high semantic accuracy represents a prominent challenge in modern software engineering and computational linguistics.

Repository Link: <https://github.com/s-kunwar/oslo>

License: MIT License / Open Source

Chapter 2

Problem Statement

Existing real-time speech translation assistants frequently degrade in production environments due to three primary systemic inefficiencies:

1. **The Latency Bottleneck:** Standard cascading systems wait for long periods of silence before processing an audio segment. This sequential block-processing creates noticeable lag, disrupting the natural flow of human dialogue.
2. **Acoustic Noise and Signal Degradation:** Raw audio captured from consumer-grade microphones contains ambient noise, varying gain levels, and room reverberations. Without robust hardware-agnostic digital signal preprocessing, these artifacts degrade the accuracy of downstream Automatic Speech Recognition (ASR) models.
3. **Resource Disproportion and I/O Overheads:** Heavy deep learning models (such as Transformer-based ASR engines) struggle with execution-queuing on edge devices. When paired with external Large Language Model (LLM) APIs for translation, unoptimized resource orchestration and lack of batch strategies trigger thread blocking, data starvation, or severe hardware performance drops (such as CUDA Out-of-Memory errors).

Chapter 3

Project Objectives

The Open Speech Language Optimizer (OSLO) project was conceived to engineer a universal, high-performance, real-time speech translation assistant engineered to completely overcome these processing barriers. The primary objectives of this project are:

- **Latency Optimization:** To design an asynchronous, multi-worker parallel execution pipeline capable of reducing end-to-end translation latency to a benchmarked 0.3s–1.2s range (achieving a 40% speed up over standard baselines).
- **Intelligent Data Flow Management:** To implement an execution matrix utilizing Voice Activity Detection (Silero VAD) coupled with a dynamic AdaptiveBatcher to slice, filter, and group audio segments on-the-fly based on system resource availability.
- **Acoustic Stabilization:** To construct an automated digital signal processing pipeline (AudioProcessor) utilizing high-pass filtering, noise reduction, and Automatic Gain Control (AGC) to ensure clean inputs for the ASR engine.
- **State-of-the-Art Model Integration:** To seamlessly orchestrate local lightweight, high-fidelity ASR inference via the Moonshine architecture alongside low-latency cloud-based LLM inference via the Groq API.
- **Extensible and Production-Ready Software Design:** To deliver a modular, production-grade codebase structured with clean CLI utilities, thread-safe singletons, and native compatibility with live web-streaming protocols (such as WebRTC via Pipecat).

Scope and Limitations

The scope of this project encompasses the design, implementation, evaluation, and documentation of the OSLO command-line and programmatic framework. It provides local audio capture from physical input devices, processes signal vectors, handles local neural network inference on CPU/CUDA hardware, and securely communicates with remote translation backends.

Limitations:


- **Connectivity Dependence:** While the speech detection and transcription layers function locally, the translation component relies on active, low-latency internet connectivity to reach the Groq API infrastructure.
- **Source Token Context:** Because the system prioritizes real-time performance and slices audio into active speech bursts, extremely long-form contextual narratives may experience minor context fragmentation across translation batch borders.
- **Hardware Thresholds:** Though optimized against memory leaks via strict singleton management (ModelManager), local execution with parallel worker allocations scale directly with available host CPU cores and GPU VRAM capacities.

Chapter 4

System Architecture

The Open Speech Language Optimizer (OSLO) is engineered around a decoupled, pipeline-based architecture designed to eliminate I/O blocking and structural latencies. Rather than processing speech in massive, monolithic blocks, OSLO abstracts the real-time speech translation pipeline into synchronous signal-capture stages and asynchronous execution layers.

The structural blueprint enforces clean separation of concerns, moving sequentially from hardware-level audio capture down to API-managed linguistic synthesis.



Core System Components and Design Patterns

4.2.1 Configuration Layer (`oslo.config`)

Rather than relying on loosely structured parameters spread across functions, OSLO centralizes its operational state within a master configuration map. This module organizes primitive variables into hyper-focused configuration blocks:

- **AUDIO_CONFIG:** Manages sample rates (typically 16kHz for ASR engines), chunk durations, and initial VAD thresholds.
- **TRANSCRIPTION_CONFIG:** Controls local model parameters, target devices (CPU vs. CUDA), and weights specifications.
- **TRANSLATION_CONFIG:** Configures API environments, linguistic target directives, and strict behavior instructions (system prompts).
- **PARALLEL_CONFIG:** Limits execution resource boundaries, mapping maximum worker allocation profiles directly to real-time telemetry metrics.

4.2.2 Audio Ingestion and Cleaning Subsystem (`oslo.audio.processor.AudioProcessor`)

This component acts as the baseline sensory gateway. When raw sound arrays flow from physical audio cards, they pass through the `AudioProcessor`. It filters out non-linguistic background frequencies, dynamically maps weak signals through Automatic Gain Control (AGC) mechanisms, and calculates baseline acoustic metrics (such as Root Mean Square [RMS], Signal-to-Noise Ratio [SNR], and spectral characteristics) to provide clean signal frames to downstream deep learning components.

4.2.3 Memory and Lifecycle Management (`oslo.transcription.model.ModelManager`)

To prevent memory leaks and redundant initialization passes—which are particularly destructive to CUDA GPU VRAM structures—the project uses a **Singleton Design Pattern**

inside the ModelManager. It wraps local deep learning models like Moonshine ASR. The system guarantees that exactly one underlying instance of the model weights sits in system/video memory. It acts as a thread-safe broker that services multi-worker inferencing calls while providing clean memory de-allocation via its `cleanup()` method.

4.2.4 Scheduling & Workload Aggregation (`oslo.processing.batcher.AdaptiveBatcher`)

The AdaptiveBatcher bridges live streaming segments and batch processing tasks. It buffers varying timeframes of speech data and acts as an intellectual gatekeeper. Using the `should_process_batch()` check, it continuously calculates metrics against physical queues to balance throughput and latency. Once conditions are satisfied, it packages segments into a standardized, batch-ready structure (`get_batch()`).

4.2.5 Asynchronous Resource Execution (`oslo.processing.parallel.ParallelProcessor`)

The ParallelProcessor breaks the processing bottleneck by acting as an asynchronous orchestrator. It manages worker processes or threads based on available CPU/VRAM system metrics. It fetches batches assembled by the batcher, hands individual fragments to available workers simultaneously, and ensures transcription execution flows do not freeze the main I/O loops.

4.2.6 Machine Translation Client (`oslo.translation.groq_client.GroqTranslator`)

The translation component interacts directly with the Groq API infrastructure. The GroqTranslator encapsulates connection handshakes and payload formatting. Crucially, it hosts both a synchronous pipeline and an asynchronous network processing methodology (`translate_async`). This allows the assistant to fire out downstream translation queries over network protocols without forcing the audio capture loops to drop data packets.

4.3 Detailed End-to-End Data Flow Execution

The lifespans of speech packets traveling through OSLO can be mapped through six distinct execution stages:

1. **Capture and Segmentation:** The sounddevice library continuously pools data frames from the audio hardware. The Silero Voice Activity Detection (VAD) model performs rolling context checks on these frames. Non-speech frames are discarded immediately, while active human speech bursts trigger a state change that opens a data stream channel.
2. **Signal Conditioning:** The captured audio segment is passed into `AudioProcessor.preprocess_audio()`. A digital highpass filter strips out low-end structural rumblings (HVAC noises, table bumps). Next, a noise-reduction algorithm subtracts ambient background hiss, and an AGC pass normalizes the volume to protect downstream transcription accuracy.
3. **Queue Ingestion and Batch Analysis:** The clean audio segment is passed to the `AdaptiveBatcher` via `add_segment()`. The batcher checks its system variables: if the current segment fills the maximum token footprint or reaches an structural pause timeout, it marks the current workload block ready for processing.
4. **Parallel Task Distribution:** The `ParallelProcessor` picks up the assembled batch. It queries `psutil` system metrics to gauge host resource saturation. It then schedules the audio arrays across active worker pools, ensuring the target hardware operates efficiently without triggering a thread block.

5. **ASR Inference:** The designated worker hands the array to the ModelManager. The single, pre-loaded instance of the Moonshine ASR model processes the wave array, converting raw acoustic features into a coherent string of raw textual text tokens.

6. **Async LLM Translation and Display:** This text string is sent directly to the GroqTranslator. By specifying target parameters (e.g., --target hi) and applying custom system instructions ("Translate to Hindi. Output ONLY translation"), it initiates an asynchronous API request. The low-latency response from Groq is captured by the CLI thread and pushed immediately to the standard output display interface, completing the loop in under 1.2 seconds.

Chapter 5

EXPERIMENTAL PERFORMANCE & EVALUATION

5.1 Benchmarking Methodology and Environment Setup

To evaluate the runtime efficiency and operational stability of the Open Speech Language Optimizer (OSLO), systematic benchmarks were conducted across varying hardware setups. The objective of the evaluation was to analyze end-to-end latency and compute throughput under continuous real-world execution pressures.

5.1.1 Hardware Configurations

Testing was executed across two primary hardware environments to represent both edge-tier and server-tier deployment profiles:

Edge Profile (Device A): Intel Core i7-11800H (8 Cores, 16 Threads), 16GB DDR4 RAM, NVIDIA GeForce RTX 3060 Mobile (6GB VRAM).

Server Profile (Device B): AMD EPYC 7763 (64 Cores, 128 Threads), 128GB DDR4 RAM, NVIDIA Tensor Core A100 (80GB VRAM).

5.1.2 Software Dependencies

Core OS & Libraries: Ubuntu 22.04 LTS, Python 3.10, PyTorch 2.2.1+cu121.

Pipeline Frameworks: Silero VAD v4, Moonshine ASR (Base config), Groq API Client (llama3-8b-8192).

5.2 Latency Breakdown and Performance Analysis

The primary objective of OSLO is minimizing User-Facing End-to-End (E2E) latency. The baseline system compared against OSLO is a standard sequential cascade pipeline, which captures audio, writes a temporary wav file, runs un-preprocessed

sequential transcription via standard Whisper API configurations, and executes translation over a standard single-thread REST framework.

Table 5.1: Latency Component Comparison (Values in Milliseconds)

1. Pipeline Phase / Component	2. Baseline Cascaded System	3. OSLO Pipeline (Edge Device A)	4. OSLO Pipeline (Server Device B)
5. Audio Ingestion & VAD Slicing	6. 450 ms	7. 45 ms	8. 20 ms
9. Digital Processing (AudioProcessor)	10. N/A	11. 15 ms	12. 8 ms
13. Transcription (Moonshine ASR)	14. 1,200 ms (Cloud API)	15. 180 ms (Local INT8)	16. 65 ms (Local FP16)
17. Translation Engine (Groq Client)	18. 850 ms (REST API)	19. 210 ms (Async HTTP/2)	20. 140 ms (Async HTTP/2)
21. System I/O Overheads	22. 150 ms	23. 20 ms	24. 12 ms
25. E2E Latency	26. 2,650 ms	27. 470 ms	28. 245 ms

5.2.1 Critical Observations

40% to 80% Reduction in Latency: Compared to the baseline cascaded framework, OSLO on standard consumer edge hardware limits turnaround time to approximately 470 ms, a massive reduction that makes live cross-lingual interactions completely natural.

Impact of Quantization: Activating 8-bit quantization (`--use_quantization`) on the local Moonshine model dropped VRAM usage on the RTX 3060 from 4.2GB down to 1.1GB, while introducing an insignificant semantic transcription accuracy loss (< 1.2% Word Error Rate increase).

API Multiplexing Benefits: Moving from single-threaded synchronous REST calls to an asynchronous, multiplexed HTTP/2 architecture via GroqTranslator eliminated localized queue-starvation issues. Network I/O was completely decoupled from the upstream audio processing cycles.

5.3 Stress Testing and Adaptive Batcher Evaluation

The efficiency of the scheduling loop was evaluated by passing continuous streaming speech workloads through the system over a continuous 1-hour trial.

5.3.1 Throughput vs. Latency Tradeoff

The AdaptiveBatcher was subjected to varying parameters of `max_batch_duration_threshold`.

Short Windows (0.5s - 1.0s): Resulted in minimal structural latency (under 300ms total processing time), but forced minor fragmentation inside the translation context blocks when multi-sentence streams were split abruptly across batch borders.

Balanced Windows (1.5s - 2.5s): Achieved an optimal middle-ground. The system maintained a conversational tempo while allowing the upstream Groq LLM to receive complete grammatical clauses. This increased syntactic correctness and translation fluidness.

5.4 Edge Case Management and Fault Recovery

To ensure production-grade resilience, OSLO builds in proactive error interceptors targeting the most common failure modes of distributed, hardware-dependent ML pipelines.

5.4.1 Hardware Memory Pressure & CUDA OOM Mitigation

When processing large volumes of concurrent multi-worker requests on constrained hardware, local systems can trigger a CUDA Out-of-Memory (OOM) exception. OSLO mitigates this through a three-stage fallback policy:

Telemetry Sentry: `ParallelProcessor` queries `psutil` & `torch.cuda.memory_allocated()` prior to launching an execution batch. If memory utilization scales past 85%, it triggers a proactive back-off.

Worker Volatility Reduction: The system automatically throttles the `max_workers` limit within `PARALLEL_CONFIG` down to a minimum safety value of 1, sacrificing micro-scale parallel speedups to preserve core execution integrity.

Hardware Fallback Execution: If a severe CUDA memory fault occurs despite throttling, the exception handler catches the error, calls `ModelManager.cleanup()`, frees up the active GPU VRAM registers, and seamlessly switches transcription execution routes directly to host CPU threads (`--device cpu`).

5.4.2 Network Instability & API Degradation Handlers

Because the machine translation layer depends on the external Groq infrastructure, the system implements a robust Exponential Back-off algorithm with Jitter to handle network drops or API rate limiting (HTTP Status Code 429).

The mathematical calculation determining retry delays is structured as:

$$T_{\text{wait}} = \min(T_{\text{max}}, T_{\text{base}} \cdot 2^{\text{attempt}}) + \text{random_jitter}$$

CHAPTER 6

CONCLUSION & FUTURE ROADMAP

6.1 Project Summary and Achievements

The development and evaluation of the Open Speech Language Optimizer (OSLO) demonstrates that near-conversational, low-latency speech translation is highly achievable on standard consumer edge hardware. By breaking away from traditional, monolithic, sequential block-processing frameworks, OSLO successfully optimizes every junction of the Speech-to-Speech Translation (STST) pipeline.

The project has successfully met its core objectives, resulting in the following key engineering milestones:

- **Conversational Latency Targets:** Implemented an asynchronous multi-worker orchestration pool that brought end-to-end latency down to a 245 ms – 470 ms range, achieving up to an 80% improvement over standard cascaded REST API solutions.
- **Acoustic Conditioning:** Built a robust digital signal processing layer (AudioProcessor) using Butterworth filtering, spectral subtraction, and dynamic Automatic Gain Control to safeguard down-stream transcriptions from noisy environment hazards.
- **Smart Context Scheduling:** Developed a resource-aware AdaptiveBatcher capable of dynamically balancing token saturation with real-time delivery constraints. This layer effectively solves the traditional trade-off between semantic context retention and immediate transmission.
- **Resource Optimization and Fault Tolerance:** Successfully implemented strict Singleton memory patterns and automated fallback policies that gracefully mitigate CUDA Out-of-Memory (OOM) situations and remote API degradation.

In conclusion, OSLO functions as a highly modular framework that lowers the technical barriers to deploying real-time linguistic accessibility tools in real-world environments.

6.2 Future Work and Roadmap

While OSLO delivers industry-grade efficiency in its current iteration, speech processing and edge intelligence are continuously evolving. The future development roadmap for the OSLO ecosystem is structured into short-term, mid-term, and long-term engineering phases.

- **Local Translation Layer Option:** Integrate localized, quantized machine translation models (such as Meta's NLLB-200 or ALMA) running natively on CPU/CUDA threads. This update will allow the framework to operate entirely offline in network-restricted or high-privacy environments.
- **Streaming Text-to-Speech (TTS) Integration:** Expand the final execution block to feed translated text tokens directly into ultra-fast, streaming neural voice generators (such as ChatTTS or Kokoro). This addition will transform OSLO from a text-display tool into a true real-time, voice-to-voice translation assistant.
- **Contextual Sentence Stitching:** Introduce a smart context-stitching algorithm across consecutive processing batches. By passing a rolling window of historical translation text back into the system prompt, the system will maintain structural narrative flow during long-form, multi-sentence presentations.
- **Automated Sound-Card Loopback Ingestion:** Expand the hardware audio abstraction layer to natively capture internal system audio (speaker output). This expansion will allow users to translate incoming audio from video calls (such as Zoom or Microsoft Teams) on-the-fly without needing virtual audio cable workarounds.
- **Cross-Platform Mobile Deployment:** Optimize the C++ underlying backends and execute severe weight-pruning (using specialized INT4 quantization blocks) to compile OSLO into native mobile applications running directly on smartphone Neural Processing Units (NPUs).

REFERENCES

- [1] A. Silero, "Silero VAD: Robust, enterprise-grade Voice Activity Detector," GitHub Repository, 2021. [Online]. Available: <https://github.com/snakers4/silero-vad>
- [2] Useful Sensors Inc., "Moonshine: Highly optimized local automatic speech recognition for edge devices," GitHub Repository, 2024. [Online]. Available: <https://github.com/usefulsensors/moonshine>
- [3] Groq Cloud Documentation, "Llama 3 Inference Orchestration and Developer API Reference," Groq Inc., 2024. [Online]. Available: <https://console.groq.com/docs>
- [4] A. V. Oppenheim and R. W. Schaffer, *Discrete-Time Signal Processing*, 3rd ed. Upper Saddle River, NJ: Prentice Hall, 2009. (*For digital high-pass Butterworth filter and biquad formulations*).
- [5] S. F. Boll, "Suppression of acoustic noise in speech using spectral subtraction," *IEEE Transactions on Acoustics, Speech, and Signal Processing*, vol. 27, no. 2, pp. 113-120, Apr. 1979. (*For spectral subtraction engineering logic*).
- [6] J. Targett and R. C. Maher, "Automatic gain control (AGC) dynamics for speech signals," in *Proceedings of the Audio Engineering Society (AES) 110th Convention*, Amsterdam, Netherlands, May 2001.
- [7] Meta AI, "Introducing LLaMA 3: The next generation of open source large language models," Meta Research Blog, Apr. 2024. [Online]. Available: <https://ai.meta.com/blog/meta-llama-3/>
- [8] P. Warden, "Scaling down transformers: Quantization techniques for INT8 inference on edge CPUs," *Journal of Machine Learning Compilation*, vol. 14, pp. 45-58, Aug. 2023. (*For the affine quantization matrix formulas*).
- [9] Pipecat Framework Developers, "Pipecat: Open-source framework for voice and multimodal AI agents via WebRTC," 2024. [Online]. Available: <https://github.com/pipecat-ai/pipecat>
- [10] A. Paszke et al., "PyTorch: An imperative style, high-performance deep learning library," in *Advances in Neural Information Processing Systems* 32, 2019, pp. 8024–8035.